

Dao Xuan Tan

AI Engineer

📍 Ho Chi Minh, Vietnam

✉ tan.daoxuan0909@gmail.com

☎ 0839 782 709

🌐 tandao0909

🌐 Đào Xuân Tân

Summary

AI engineer who implements and integrates AI/ML solutions from requirement to delivery, with hands-on breadth across machine learning, deep learning, NLP, and computer vision on a strong Python foundation. Builds and serves LLM systems, fine-tunes models, and evaluates outputs for quality; researches and proposes practical solutions, and keen to keep learning new technologies.

Experience

AcamarAI, AI Engineer

Python, PyTorch, LLM/VLM, fine-tuning, NLP, RAG, vLLM, FastAPI

Ho Chi Minh

June 2025 – Jan 2026

8 months

- Delivered **3 of 14** core AI features end to end, from requirement to production, for **Fortune 500** beta users, turning product requirements into working AI/ML workflows
- Built **NLP** workflows for text **classification** and **entity resolution**, combining deterministic rules with LLMs for reliable, production data processing
- Applied **computer vision** through **ColPali** vision-RAG and **VLM** document understanding over document images, plus **BGE** embedding and ingestion pipelines
- Fine-tuned task-specific models with **LoRA/PyTorch** and prepared datasets to improve accuracy for target tasks (**deep learning**)
- Built **evaluation** harnesses (**recall@k**, **LLM-as-judge**, human-in-the-loop review) to measure and improve model output quality
- Served self-hosted **Qwen** models with **vLLM** behind an OpenAI-compatible API, shipped via async **FastAPI**

Pi Associates, Software Engineer

Python, pandas, PostgreSQL, Kafka, Redis, Docker, Kubernetes, GCP

Ho Chi Minh

Jan 2026 – present

7 months

- Main developer of an internal **Python** data and backtesting platform, used by about 20 researchers and engineers running **tens of thousands** of runs a day
- Diagnosed and fixed a critical **memory leak**, cutting peak memory from **200+ GB to under 20 GB**, and cut runtime from **minutes to under 10 seconds** on **1-5 GB** of data per run through parallelism, vectorization, and algorithmic improvements
- Built the **pandas** backtest pipeline, **PostgreSQL** schema and queries, **Redis** caching, and a **Kafka** producer feeding downstream consumers
- Ran jobs as Dockerized **Kubernetes** workloads on **GCP**, with **GitHub Actions** CI/CD and Grafana monitoring
- Improved reliability with stronger error handling and input validation, reducing failed runs and manual debugging

Skills

Programming: Python (async, typing, packaging), SQL, problem-solving, algorithms and data structures

Machine Learning & Deep Learning: PyTorch, TensorFlow, LoRA fine-tuning, dataset preparation, model evaluation, embeddings

NLP: text classification, entity resolution, retrieval (RAG), text generation, prompt engineering

Computer Vision: vision-language models (VLM), ColPali document understanding, image-based retrieval

LLM & Agents: vLLM self-hosted serving, OpenAI-compatible APIs, RAG, tool-calling agents, LLM-as-judge evaluation

Backend & Data: FastAPI, REST APIs, PostgreSQL, Redis, Kafka, pandas, data pipelines

Cloud & DevOps: GCP, Docker, Kubernetes, GitHub Actions CI/CD, Grafana

Languages: Vietnamese (Native), English (IELTS 6.5)

Education

Ho Chi Minh University of Science, Mathematics and Computer Science, Data Science track

- GPA: 8.74/10.0 (3.74/4.0)

Ho Chi Minh City
Oct 2022 – Sept 2026