

Dao Xuan Tan

Software Engineer (AI-Native)

📍 Ho Chi Minh, Vietnam

✉ tan.daoxuan0909@gmail.com

☎ 0839 782 709

🔗 tandao0909

🌐 Đào Xuân Tân

Summary

Software engineer who builds, tests, and reviews with agentic AI coding tools (Claude Code, Cursor, Codex) every day and verifies their output before it ships. Backed by production **Python**, large-scale data pipelines, and applied ML (classification, embeddings, evaluation), with a habit of chasing quality problems down to their root cause.

Experience

Pi Associates, Software Engineer

Ho Chi Minh

Python, PostgreSQL, SQL, Kafka, Redis, Docker, Kubernetes, GCP, GitHub Actions, Grafana

Jan 2026 – present

7 months

- Main developer responsible for an internal **Python** data and backtesting package, used by about 20 researchers and engineers running **tens of thousands** of runs a day
- Build, test, and review features daily with agentic AI coding tools (**Claude Code**, **Cursor**, **Codex**), verifying every change against tests and real data before merging
- Built and improved data pipelines over **PostgreSQL** (schema design and SQL queries), **Redis** caching, and a **Kafka** producer feeding signals to downstream consumers
- Traced and fixed a silent production memory leak no one had caught, cutting peak memory from **200+ GB to under 20 GB** by finding the root cause and shipping the fix
- Cut runtime from **minutes to under 10 seconds** on **1-5 GB** of data per run through parallelism, vectorization, and algorithmic improvements
- Ran jobs as Dockerized **Kubernetes** workloads on **GCP** and **Linux**, with **GitHub Actions** CI/CD and **Grafana** production monitoring
- Improved reliability with stronger error handling and input validation, cutting failed runs and manual debugging

AcomarAI, AI Engineer

Ho Chi Minh

Python, LLM, classification, entity resolution, embeddings, evaluation, RAG, vLLM

June 2025 – Jan 2026

8 months

- Built **classification** and **entity-resolution** workflows that combine deterministic rules with **LLMs** for reliable data processing in production
- Built **evaluation harnesses**, automated output scoring on metrics like **recall@k** and **LLM-as-judge** plus human-in-the-loop review, to measure and improve detection quality
- Built retrieval and **RAG** pipelines over structured and unstructured data, parsing and chunking documents and indexing them with **BGE** embeddings
- Deployed self-hosted **LLM** inference (**vLLM**) behind an OpenAI-compatible API for internal AI features
- Owned **3** production AI features end to end, out of **14** core features, for Fortune 500 beta users

Skills

AI-Native Development: Claude Code, Cursor, Copilot and Codex agentic coding and review, prompt engineering, LLM APIs (OpenAI, Claude, Gemini), verifying and testing AI output

ML & Data Intelligence: classification, entity resolution, embeddings (BGE), RAG, evaluation metrics (recall@k, LLM-as-judge)

Languages & CS: Python (asyncio, typing, packaging with uv), C++ (coursework and competitive programming), SQL, data structures and algorithms

Backend & Data Pipelines: FastAPI, REST and async APIs, PostgreSQL, Redis, Kafka, Pandas, large-scale data pipelines, parsing structured and unstructured data

Infra & DevOps: Docker, Kubernetes, GCP, Linux, GitHub Actions CI/CD, Grafana, Git, Pytest

Languages: Vietnamese (Native), English (IELTS 6.5)

Education

Ho Chi Minh University of Science, Mathematics and Computer Science, Data Science track

Ho Chi Minh City

- GPA: 8.74/10.0 (3.74/4.0)

Oct 2022 – Sept 2026