

Dao Xuan Tan

AI Solution Engineer

📍 Ho Chi Minh, Vietnam

✉ tan.daoxuan0909@gmail.com

☎ 0839 782 709

🌐 tandao0909

🌐 Đào Xuân Tân

Summary

AI engineer building production LLM systems for legal, IP, and compliance workflows, with hands-on experience in RAG pipelines, vector search, tool-calling agents, and backend API integration. Ship schema-validated, human-reviewed AI automation using Python/FastAPI, OpenAI-compatible APIs, self-hosted models, and reliability-focused engineering practices.

Experience

AcamarAI, AI Engineer

Ho Chi Minh

Legal/IP LLM systems, RAG, Qdrant, tool-calling agents, LangChain, Instructor, vLLM, FastAPI

June 2025 – Jan 2026

8 months

- Owned end-to-end **LLM patent drafting and classification** system, one of **3 of 14** core features for **Fortune 500** beta client; generated patent claim and description text, classified technical documents, and turned legal and **IP** requirements into production AI workflow
- Enforced reliable, well-formed legal output via **schema-validated** generation (**Instructor/Pydantic**); validated quality through **patent and domain-expert human-in-the-loop review** and **LLM-as-judge** evaluation harnesses
- Built production **RAG and retrieval** over **millions of vectors** in self-hosted **Qdrant**, combining **hybrid** dense and sparse (BM25) search with **reranking** model that outperformed pure vector search on **recall@k** for prior-art-style retrieval across legal and technical corpora
- Developed **agentic** LLM workflows on **OpenAI tool calling** and **LangChain** with multi-step planning and tools for text processing and internal-DB retrieval; hardened reliability against retries and malformed tool arguments, plus ingestion pipelines using **BGE** embeddings and **ColPali** vision RAG over document images
- Served self-hosted **Qwen** models with **vLLM** on single A100 80GB behind **OpenAI-compatible API** shipped features via **async FastAPI** with **SSE** streaming

Pi Associates, Software Developer (Quant Finance)

Ho Chi Minh

Python, FastAPI, PostgreSQL, Kafka, Redis, Docker, Kubernetes, GCP

Jan 2026 – present

6 months

- Main developer of internal **Python** backtesting platform (FastAPI, PostgreSQL, Kafka, Redis) serving about 20 researchers running **tens of thousands** of backtests daily
- Diagnosed and fixed critical **memory leak**, cutting peak memory from **200+ GB to under 20 GB** reduced runtime from **minutes to under 10 seconds** on **1-5 GB** per run through **parallelism** and vectorization
- Ran backtests as Dockerized **Kubernetes** jobs across **GCP** and on-premise, with **GitHub Actions** CI/CD and **Grafana** monitoring

Skills

LLM & Agents: RAG, OpenAI tool/function calling, multi-step planning agents, LangChain, Instructor, schema-validated outputs, prompt engineering, LLM-as-judge evaluation, AI-assisted dev tools (Codex, Claude Code, Cursor)

Retrieval & Vector DBs: Qdrant, hybrid dense + sparse (BM25) retrieval, reranking models, BGE and ColPali embeddings, chunking strategies, recall@k

LLM Serving & ML: vLLM, self-hosted Qwen, OpenAI-compatible APIs, LoRA fine-tuning, PyTorch, VLM and document understanding

Backend & APIs: Python, FastAPI, asyncio, SSE streaming, REST APIs, Pydantic validation, PostgreSQL, Redis, Kafka

Infra & DevOps: Docker, Kubernetes, GCP, on-premise, GitHub Actions CI/CD, Grafana

Legal/IP Domain: patent drafting and classification, IP-document RAG, compliance-oriented evaluation and human-in-the-loop review

Languages: Vietnamese (Native), English (IELTS 6.5)

Education

Ho Chi Minh University of Science, Mathematics & Computer Science (Data Science)

Ho Chi Minh City

- GPA: 8.74/10.0 (3.74/4.0)

Oct 2022 – Sept 2026